

# multilinearregression

August 24, 2023

## 1 Multiple Linear Regression

- Linear Regression is used for predicting the relationship between the independent and dependent variable. It aims to predict the good fit line that describes the relationship by minimizing the mean\_squared\_error of the predicted value and the actual value.
- In multiple linear regression, we use multiple/more than one independent variable/feature to predict the target variable.

- Slope-intercept formulae,  $Y=B_0+B_1X_1+B_2X_2+\dots+B_iX_i$

where,  $y$ = target variable( Dependent variable)  $B_0$ = intercept  $B_1$ =slope/coefficient  $X_1, X_2, X_3, \dots, X_i$ =Predictive variables(Independent variable)

- Linear regression aims to find the slope value and intercepts value(in case of multiple linear model- $B_1, B_2, B_3, \dots$ ) to predict the target feature.
- Have used the Car Price Estimation data set from kaggle.com; where, Independent variable(X)- 'car\_name', 'fuel\_type', 'no\_cylinder', 'seating\_capacity', 'transmission\_type', 'body\_type', 'start' and target variable(Y)- "ending\_price"

```
[1]: # Importing necessary libraries
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: # Onboarding data onto colab
```

```
from google.colab import files
rawdata=files.upload()
```

<IPython.core.display.HTML object>

Saving cars.csv to cars.csv

```
[3]: # Converting to dataframe
```

```
df=pd.read_csv('cars.csv')
```

df

[3]:

	car_name	reviews_count	fuel_type	\
0	Maruti Alto K10	51	Petrol	
1	Maruti Brezza	86	Petrol	
2	Mahindra Thar	242	Diesel	
3	Mahindra XUV700	313	Diesel	
4	Mahindra Scorpio-N	107	Diesel	
..	...	...	...	
198	Mercedes-Benz AMG A 45 S	35	Petrol	
199	BMW 3 Series Gran Limousine	3	Petrol	
200	MG Hector Plus	2	Diesel	
201	Audi RS Q8	9	Petrol	
202	Maruti Alto 800 tour	4	Petrol	

	engine_displacement	no_cylinder	seating_capacity	transmission_type	\
0	998	3	5.0	Automatic	
1	1462	4	5.0	Automatic	
2	2184	4	4.0	Automatic	
3	2198	4	7.0	Automatic	
4	2198	4	7.0	Automatic	
..	...	...	...	...	
198	1991	4	5.0	Automatic	
199	1998	4	5.0	Automatic	
200	1956	4	7.0	Manual	
201	3998	8	5.0	Automatic	
202	796	3	5.0	Manual	

	fuel_tank_capacity	body_type	rating	starting_price	ending_price	\
0	27.0	Hatchback	4.5	399000	583000	
1	48.0	SUV	4.5	799000	1396000	
2	57.0	SUV	4.5	1353000	1603000	
3	60.0	SUV	4.5	1318000	2458000	
4	57.0	SUV	4.5	1199000	2390000	
..	...	...	...	...	...	
198	0.0	Hatchback	4.5	659000	999000	
199	59.0	Sedan	4.5	1041000	1041000	
200	60.0	SUV	4.5	1615000	2075000	
201	85.0	SUV	3.5	21700000	21700000	
202	35.0	Hatchback	4.5	391000	397000	

	max_torque_nm	max_torque_rpm	max_power_bhp	max_power_rp
0	89.0	3500	65.71	5500
1	136.8	4400	101.65	6000
2	300.0	2800	130.00	3750
3	450.0	2800	182.38	3500
4	400.0	2750	172.45	3500

..	...	...	...	...
198	500.0	5250	415.71	6750
199	400.0	4400	254.79	5000
200	350.0	2500	167.67	3750
201	800.0	4500	591.39	6000
202	69.0	3500	47.33	6000

[203 rows x 16 columns]

```
[4]: # Shallow copy
df_copy=df.copy()
```

## 2 Exploratory Data Analysis

```
[5]: df.head()
```

```
[5]:
```

	car_name	reviews_count	fuel_type	engine_displacement	\
0	Maruti Alto K10	51	Petrol	998	
1	Maruti Brezza	86	Petrol	1462	
2	Mahindra Thar	242	Diesel	2184	
3	Mahindra XUV700	313	Diesel	2198	
4	Mahindra Scorpio-N	107	Diesel	2198	

	no_cylinder	seating_capacity	transmission_type	fuel_tank_capacity	\
0	3	5.0	Automatic	27.0	
1	4	5.0	Automatic	48.0	
2	4	4.0	Automatic	57.0	
3	4	7.0	Automatic	60.0	
4	4	7.0	Automatic	57.0	

	body_type	rating	starting_price	ending_price	max_torque_nm	\
0	Hatchback	4.5	399000	583000	89.0	
1	SUV	4.5	799000	1396000	136.8	
2	SUV	4.5	1353000	1603000	300.0	
3	SUV	4.5	1318000	2458000	450.0	
4	SUV	4.5	1199000	2390000	400.0	

	max_torque_rpm	max_power_bhp	max_power_rp
0	3500	65.71	5500
1	4400	101.65	6000
2	2800	130.00	3750
3	2800	182.38	3500
4	2750	172.45	3500

```
[6]: df.columns
```

```
[6]: Index(['car_name', 'reviews_count', 'fuel_type', 'engine_displacement',
         'no_cylinder', 'seating_capacity', 'transmission_type',
         'fuel_tank_capacity', 'body_type', 'rating', 'starting_price',
         'ending_price', 'max_torque_nm', 'max_torque_rpm', 'max_power_bhp',
         'max_power_rp'],
        dtype='object')
```

```
[7]: # Feature Engineering

df=df.
↳drop(['engine_displacement','fuel_tank_capacity','rating','max_torque_nm',
↳'max_torque_rpm',
↳'max_power_bhp','reviews_count','max_power_rp'],axis='columns')
```

```
[8]: df
```

```
[8]:
```

	car_name	fuel_type	no_cylinder	seating_capacity	\
0	Maruti Alto K10	Petrol	3	5.0	
1	Maruti Brezza	Petrol	4	5.0	
2	Mahindra Thar	Diesel	4	4.0	
3	Mahindra XUV700	Diesel	4	7.0	
4	Mahindra Scorpio-N	Diesel	4	7.0	
..	...	...	...	...	
198	Mercedes-Benz AMG A 45 S	Petrol	4	5.0	
199	BMW 3 Series Gran Limousine	Petrol	4	5.0	
200	MG Hector Plus	Diesel	4	7.0	
201	Audi RS Q8	Petrol	8	5.0	
202	Maruti Alto 800 tour	Petrol	3	5.0	

	transmission_type	body_type	starting_price	ending_price
0	Automatic	Hatchback	399000	583000
1	Automatic	SUV	799000	1396000
2	Automatic	SUV	1353000	1603000
3	Automatic	SUV	1318000	2458000
4	Automatic	SUV	1199000	2390000
..	...	...	...	...
198	Automatic	Hatchback	659000	999000
199	Automatic	Sedan	1041000	1041000
200	Manual	SUV	1615000	2075000
201	Automatic	SUV	21700000	21700000
202	Manual	Hatchback	391000	397000

```
[203 rows x 8 columns]
```

```
[9]: # Technical enquiry
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 203 entries, 0 to 202
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_name              203 non-null   object
1   fuel_type             203 non-null   object
2   no_cylinder           203 non-null   int64
3   seating_capacity      202 non-null   float64
4   transmission_type     203 non-null   object
5   body_type             203 non-null   object
6   starting_price        203 non-null   int64
7   ending_price          203 non-null   int64
dtypes: float64(1), int64(3), object(4)
memory usage: 12.8+ KB
```

```
[10]: # Statistical enquiry
```


```
df.describe()
```

```
[10]:
```

	no_cylinder	seating_capacity	starting_price	ending_price
count	203.000000	202.000000	2.030000e+02	2.030000e+02
mean	4.709360	5.014851	9.443640e+06	1.112005e+07
std	2.538664	1.161050	1.357035e+07	1.551746e+07
min	0.000000	2.000000	3.390000e+05	3.610000e+05
25%	4.000000	5.000000	9.455000e+05	1.407500e+06
50%	4.000000	5.000000	4.312000e+06	4.600000e+06
75%	6.000000	5.000000	1.160000e+07	1.575000e+07
max	12.000000	8.000000	7.060000e+07	9.000000e+07

```
[11]: sns.
```


```
pairplot(df,x_vars=['car_name','fuel_type','no_cylinder','seating_capacity','transmission_t',
plt.show()
```



```
[12]: sns.heatmap(df.corr())
plt.show()
```

<ipython-input-12-88edb43bf50b>:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only

```
to silence this warning.  
sns.heatmap(df.corr())
```



```
[13]: # Finding Nan values
```

```
df.isnull().sum()
```

```
[13]: car_name          0  
fuel_type          0  
no_cylinder        0  
seating_capacity    1  
transmission_type  0  
body_type          0  
starting_price     0  
ending_price       0  
dtype: int64
```


```
[14]: df.dropna(inplace=True)
```

```
[16]: # Finding duplicates if any,  
df.duplicated().sum()
```

```
[16]: 0
```

### 3 Multivariate Analysis

```
[45]: sns.  
    > lmpplot(x='body_type',y='seating_capacity',data=df,hue='transmission_type',markers=['o','*'],  
    plt.show()
```



### 4 Data encoding

```
[17]: from sklearn.preprocessing import LabelEncoder  
le=LabelEncoder()
```

```
df['car_name']=le.fit_transform(df['car_name'])
df['fuel_type']=le.fit_transform(df['fuel_type'])
df['transmission_type']=le.fit_transform(df['transmission_type'])
df['body_type']=le.fit_transform(df['body_type'])
```

```
[18]: df
```

```
[18]:
```

	car_name	fuel_type	no_cylinder	seating_capacity	transmission_type	\
0	106	3	3	5.0	0	
1	108	3	4	5.0	0	
2	101	1	4	4.0	0	
3	103	1	4	7.0	0	
4	100	1	4	7.0	0	
..	...	...	...	...	...	
198	127	3	4	5.0	0	
199	12	3	4	5.0	0	
200	89	1	4	7.0	2	
201	7	3	8	5.0	0	
202	105	3	3	5.0	2	

	body_type	starting_price	ending_price
0	2	399000	583000
1	8	799000	1396000
2	8	1353000	1603000
3	8	1318000	2458000
4	8	1199000	2390000
..	...	...	...
198	2	659000	999000
199	9	1041000	1041000
200	8	1615000	2075000
201	8	21700000	21700000
202	2	391000	397000

[202 rows x 8 columns]

```
[19]: X=df.iloc[:, :-1]
```

```
[20]: Y=df.iloc[:, -1]
```

## 5 Normalization

```
[21]: from sklearn.preprocessing import StandardScaler

sc=StandardScaler()
X_sc=sc.fit_transform(X)
```



## 6 Train test split

```
[24]: from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test=train_test_split(X_sc,Y,test_size=0.
↪3,random_state=51)
```

## 7 Model Training

```
[25]: from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
[25]: LinearRegression()
```

```
[26]: # Predicting the output

y_predict=lr.predict(x_test)
```

```
[27]: y_predict
```

```
[27]: array([ 5450734.87876823, 28883042.20779844,  4830941.34009862,
        1087634.03203458, 28701197.29354954,  1499695.85680375,
        2310424.10412433, 14674775.70844833, 17934581.96638384,
        27757791.07494777, 10128661.25953963, 38009168.96838657,
         831654.32959579,  685638.64178005,  486334.52338482,
        2713701.43982593, 1324169.33806021,  588217.97080897,
        3214435.02888541, -279229.20308003,  4786805.61487156,
         914268.63947213, 18452982.23247068, 24314601.25868671,
        21267974.07534245, 75207031.85186312,  9396687.67137627,
        15048608.02795183, 10202771.12999566,  1890689.88796972,
        2257427.57043463, 11301089.0161855 ,  9011895.71803283,
        5572125.85625006, 16294798.48283658, 15184742.40018022,
        1103229.21737047,  9478188.42698938,  6571878.88623431,
        19450456.39461272, 18541788.49635421,  1655876.11357459,
        7897797.73433908,  3540427.63245242,  5159408.95108046,
        4566098.34026237,  2279503.69999649,  8936771.034402 ,
        4006078.11742655,  1751211.43100507,  1673271.3206316 ,
        67488597.67475884, 15746785.31150446,  1774231.44549256,
        4283497.83869911,  973177.57010573,  1315179.70276314,
        1070705.19787242,  2778785.12087851,  1861267.76001905,
        21042372.77040582])
```

```
[28]: y_test
```

```
[28]: 68      4435000
      121     27100000
      90      3188000
       7       949000
      120     25500000
      ...
      19       918000
      32      1079000
      27      1549000
      51      1949000
      124     22200000
      Name: ending_price, Length: 61, dtype: int64
```

## 8 Model Evaluation

```
[29]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
MAE=mean_squared_error(y_test,y_predict)
MSE=mean_absolute_error(y_test,y_predict)
RMSE=np.sqrt(MSE)
print("MAE",MAE)
print("MSE",MSE)
print("RMSE",RMSE)
```

```
MAE 20002026248420.62
MSE 2237738.5409558667
RMSE 1495.907263487903
```

```
[30]: R2=r2_score(y_test,y_predict)
      print("R2",R2)
```

```
R2 0.9276182330996255
```

R2 score reveals that the model accuracy is 0.9 (i.e) 90% of the data fits in the good fit line/regression line.